


# Classification Error Metrics



**METIS**

**datascope**

%54 Democrats, %46 republicans

Classify:

Predict party using votes

%54 Democrats, %46 republicans

Classify:

Predict party using votes

Model performance:

"How many times did I get it right?"

%54 Democrats, %46 republicans

Classify:

Predict party using votes

Model performance:

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

%54 Democrats, %46 republicans

Classify:

Predict party using votes

Model performance:

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

95% accuracy: Good job!

%1 have leukemia, %99 are healthy

Classify:

Predict leukemia using health records and tests

%1 have leukemia, %99 are healthy

Classify:

Predict leukemia using health records and tests

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

%1 have leukemia, %99 are healthy

Classify:

Predict leukemia using health records and tests

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

Imagine the stupidest predictor:

"Always predict healthy"



%1 have leukemia, %99 are healthy

Classify:

Predict leukemia using health records and tests

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

Imagine the stupidest predictor:

"Always predict healthy"

What will its accuracy be?

%1 have leukemia, %99 are healthy

Classify:

Predict leukemia using health records and tests

"How many times did I get it right?"

**Accuracy.** % correct prediction of all predictions

Imagine the stupidest predictor:

"Always predict healthy"

What will its accuracy be?

It will be correct 99% of the time!

You won't catch any sick people. **Useless.**

# Confusion Matrix

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

# Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)	Accuracy
Spam (Actual)	27	6	81.81
Non-Spam (Actual)	10	57	85.07
Overall Accuracy			83.44

# Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)	Accuracy
Spam (Actual)	27	6	81.81
Non-Spam (Actual)	10	57	85.07
Overall Accuracy			83.44

Sensitivity

Specificity

Accuracy

# Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)	Accuracy
Spam (Actual)	0	10	0.0
Non- Spam (Actual)	0	990	100.00
Overall Accuracy			99

Sensitivity

Specificity

Accuracy

# Confusion Matrix

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

$$\text{Sensitivity} = TP / (TP + FN)$$

$$\text{Specificity} = TN / (TN + FP)$$

$$\text{Accuracy} = (TP + TN) / (TP + TN + FP + FN)$$

# Confusion Matrix

	P' (Predicted)	N' (Predicted)
P (Actual)	True Positive	False Negative
N (Actual)	False Positive	True Negative

Sensitivity =  $TP / (TP + FN)$

TRUE POSITIVE RATE

Specificity =  $TN / (TN + FP)$

TRUE NEGATIVE RATE

Accuracy =  $(TP + TN) / (TP + TN + FP + FN)$



# Precision and recall

# Precision and recall

**Precision.** Out of all cases I predicted as positive, how many times was I right?  
(% times I was right when I told somebody they had leukemia)

**Recall.** Out of all the (few) positive cases, how many did I find?  
(% of actual leukemia patients I could catch with my classifier)

# Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)	Accuracy
Spam (Actual)	27	6	81.81
Non-Spam (Actual)	10	57	85.07
Overall Accuracy			83.44

**Precision** =  $27 / 37 = 73.0\%$

**Recall** =  $27 / 33 = 81.8\%$

# Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)	Accuracy
Spam (Actual)	0	10	0.0
Non-Spam (Actual)	0	990	100.00
Overall Accuracy			99

**Precision** =  $0 / 0$  Undefined!

**Recall** =  $2 / 10 = 20\%$

# Confusion Matrix

	P' (Predicted)	n' (Predicted)
P (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

**Precision** =  $TP / (TP + FP)$

**Recall** =  $TP / (TP + FN)$

# Confusion Matrix

	P' (Predicted)	n' (Predicted)
P (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Focusing on a  
single class  
(positive: the  
one with small  
prevalence) in  
skewed cases

**Precision** =  $TP / (TP + FP)$

**Recall** =  $TP / (TP + FN)$

# Confusion Matrix

	P' (Predicted)	n' (Predicted)
P (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Focusing on a  
single class  
(positive: the  
one with small  
prevalence) in  
skewed cases


**Precision** =  $TP / (TP + FP)$

**Recall** =  $TP / (TP + FN)$

**F1** = Their harmonic mean

$$F1 = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$


# Logistic regression




$$y_{\beta}(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$



# Logistic regression




# Logistic regression

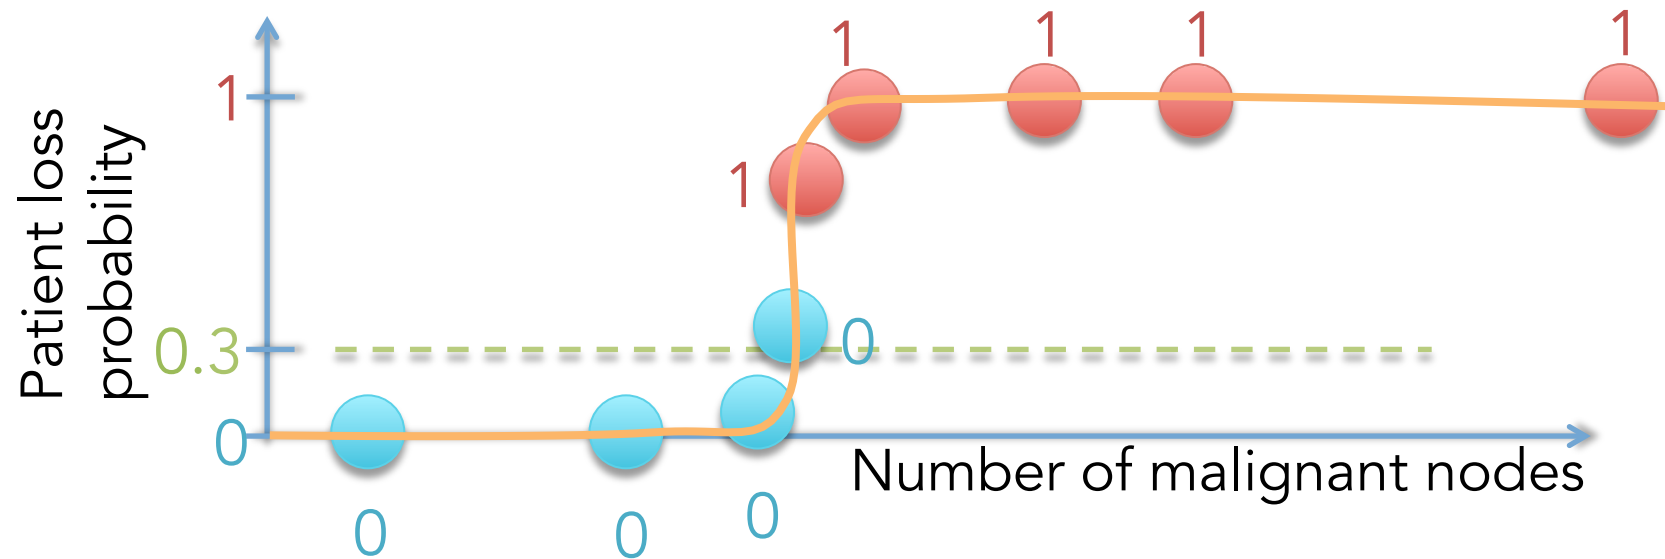


Higher threshold: Only call it positive if absolutely sure  
To make most positive calls right, give up on catching all  
lower recall, higher precision  
lower True Positive Rate, lower False Positive Rate


# Logistic regression



# Logistic regression



# Logistic regression



Lower threshold: Ease criteria for calling it positive  
Can catch more positives by accepting edge cases  
higher recall, lower precision  
higher True Positive Rate, higher False Positive Rate

# How to tune precision and recall?

**Google.**

**Forget Recall.** There are millions of relevant search results. Catching all of them is not important: Most people only look at the first page of results.

**Precision way more important.** If one of those ten links in the first page is not relevant to the search, you'll lose customers.

# How to tune precision and recall?

## HIV Tests.

**Recall very very important.** We want to catch all sick people. Imagine an HIV+ person taking the test and being told they are HIV-. Worst outcome. Nope.

**Precision also important.** Telling someone that they are HIV+ when they are not is a devastating mistake.

Not as bad as letting HIV+ slip through, but we still cannot turn the knob all the way to recall without a care about this, either.

# How to tune precision and recall?

## United States Legal System.

**Higher Precision at the cost of Recall.** Burden of proof. Reasonable doubt = Acquittal. These mean that the “threshold” for classifying someone with the ‘criminal’ label is set very high.

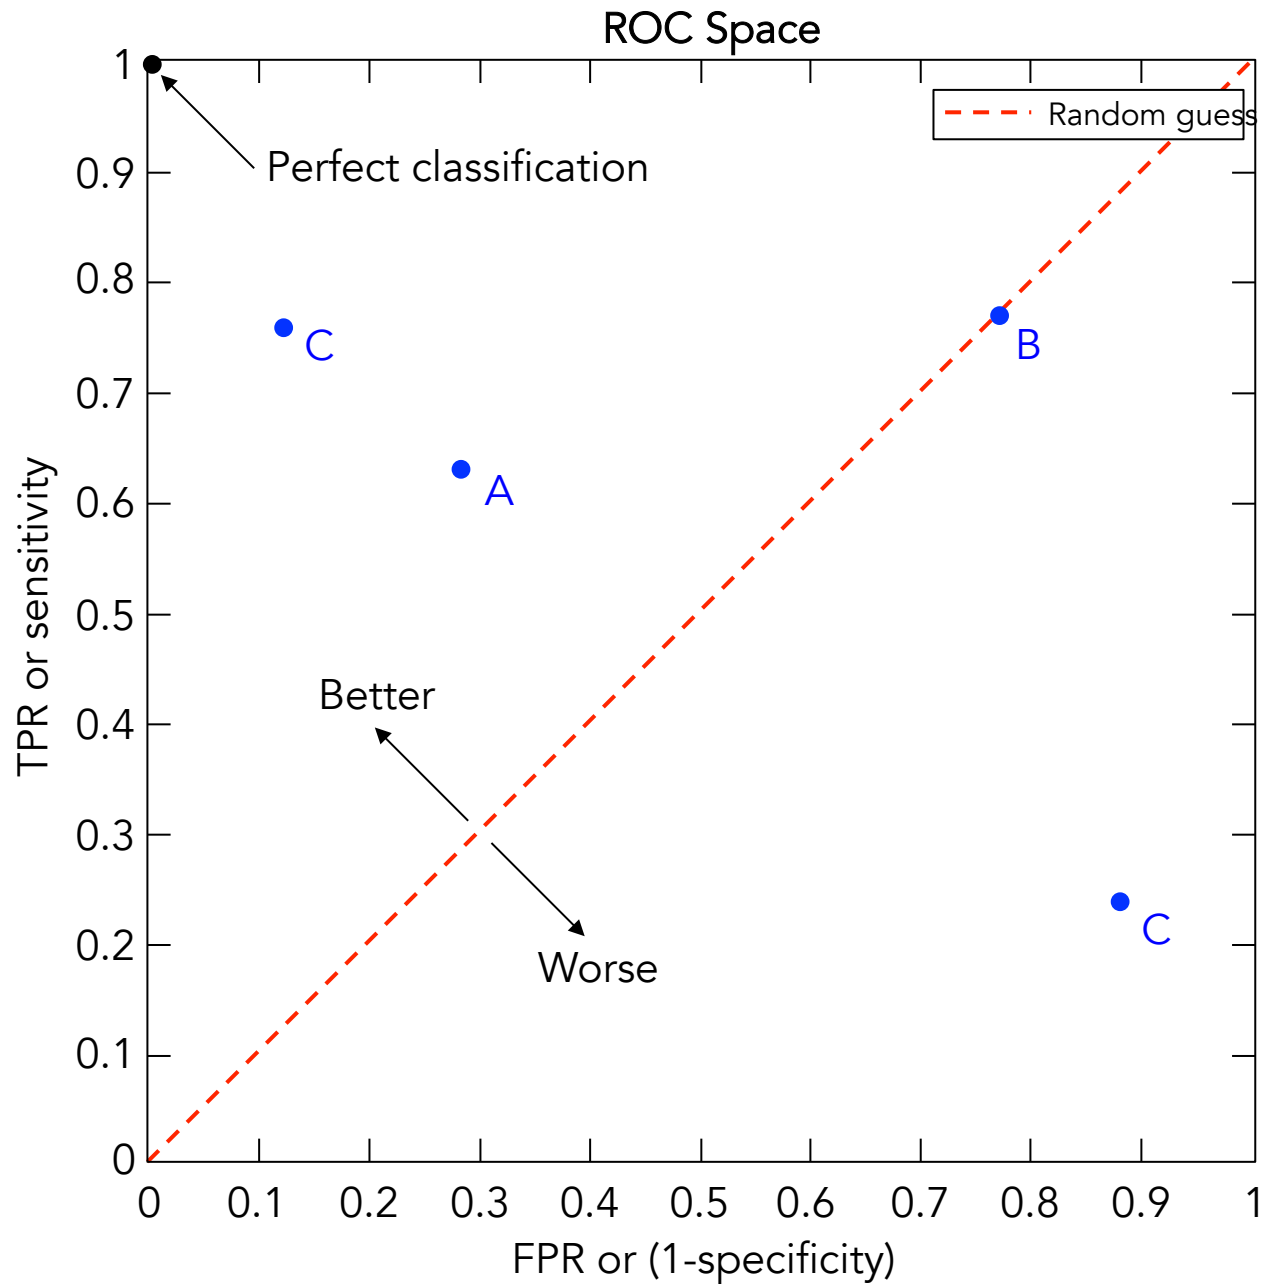
**High precision** means making sure everyone sent to jail is indeed guilty.


**High recall** means making sure all criminals go to jail. The legal system is ready to pay the cost of letting some criminals go to ensure no innocents go to jail.




Each threshold is a different model

Plot their True Positive Rate & False Positive Rate







# Receiver Operating Characteristic



Lower threshold: Ease criteria for calling it positive  
Can catch more positives by accepting edge cases  
higher recall, lower precision  
higher True Positive Rate, higher False Positive Rate




Higher threshold: Only call it positive if absolutely sure  
To make most positive calls right, give up on  
lower recall, higher precision  
lower True Positive Rate, lower False Positive Rate



## **Area under curve (AUC)**

An evaluation of a classification algorithm  
(including all possible thresholds)



# from sklearn.metrics import .....

## Classification metrics

See the [Classification metrics](#) section of the user guide for further details.

<code>metrics.accuracy_score(y_true, y_pred[, ...])</code>	Accuracy classification score.
<code>metrics.auc(x, y[, reorder])</code>	Compute Area Under the Curve (AUC) using the trapezoidal rule
<code>metrics.average_precision_score(y_true, y_score)</code>	Compute average precision (AP) from prediction scores
<code>metrics.classification_report(y_true, y_pred)</code>	Build a text report showing the main classification metrics
<code>metrics.confusion_matrix(y_true, y_pred[, ...])</code>	Compute confusion matrix to evaluate the accuracy of a classification
<code>metrics.f1_score(y_true, y_pred[, labels, ...])</code>	Compute the F1 score, also known as balanced F-score or F-measure
<code>metrics.fbeta_score(y_true, y_pred, beta[, ...])</code>	Compute the F-beta score
<code>metrics.hamming_loss(y_true, y_pred[, classes])</code>	Compute the average Hamming loss.
<code>metrics.hinge_loss(y_true, pred_decision[, ...])</code>	Average hinge loss (non-regularized)
<code>metrics.jaccard_similarity_score(y_true, y_pred)</code>	Jaccard similarity coefficient score
<code>metrics.log_loss(y_true, y_pred[, eps, ...])</code>	Log loss, aka logistic loss or cross-entropy loss.
<code>metrics.matthews_corrcoef(y_true, y_pred)</code>	Compute the Matthews correlation coefficient (MCC) for binary classes
<code>metrics.precision_recall_curve(y_true, ...)</code>	Compute precision-recall pairs for different probability thresholds
<code>metrics.precision_recall_fscore_support(...)</code>	Compute precision, recall, F-measure and support for each class
<code>metrics.precision_score(y_true, y_pred[, ...])</code>	Compute the precision
<code>metrics.recall_score(y_true, y_pred[, ...])</code>	Compute the recall
<code>metrics.roc_auc_score(y_true, y_score[, ...])</code>	Compute Area Under the Curve (AUC) from prediction scores
<code>metrics.roc_curve(y_true, y_score[, ...])</code>	Compute Receiver operating characteristic (ROC)
<code>metrics.zero_one_loss(y_true, y_pred[, ...])</code>	Zero-one classification loss.



Always remember,

Fit the model to a **training set**,

Calculate performance

(accuracy, precision, recall, f1, AUC, etc.)

on a **test set**

or (better) on a k-fold **cross validation** scheme

```
from sklearn.cross_validation import cross_val_score
```

```
KNN = KNeighborsNeighbors( 10 )
```

```
cross_val_score( KNN, scoring="recall" )
```